

Toolbox “DEEPGLM”

Type	Toolbox
Title	Bayesian Deep Net Generalized Linear Model (deepGLM)
Version	1.0
Released	May 2018
Description	Matlab Toolbox for paper Bayesian Deep Net GLM and GLMM with Natural Gradient Factor-Gaussian Variational Approximation (DeepGLM, 2018) by Minh-Ngoc Tran, Nghia Nguyen, David J. Nott and Robert Kohn. DeepGLM is a flexible version of Generalized Linear Model where Deep Feedforward Network is used to automatically choose transformations for the raw covariates.
Language	Python 3.6
Author	Nghia Nguyen (nghia.nguyen@sydney.edu.au) Minh-Ngoc Tran (minh-ngoc.tran@sydney.edu.au)

Contents

Download and Install.....	2
Preparing Data.....	2
deepGLMfit.....	3
deepGLMpredict.....	6
deepGLMplot.....	7
Practical Recommendation.....	8

Download and Install

Download and use deepGLM Python package

- Step 1: Clone or download the deepGLM.pyc file from GitHub link:
<https://github.com/VBayesLab/deepGLM>
- Step 2: Copy the deepGLM.pyc to your project folder

Preparing Data

Preparing data for deepGLM

To fitting a deepGLM model, prepare your data in the form that is accepted by fitting function `deepGLMfit`. The data for training and testing must be divided into two parts, a design matrix X for predictor variables and a column y for response variable.

- Predictor variables X , specified as an n -by- p matrix, where n is the number of observations and p is the number of predictor variables. Each column of X represents one variable, and each row represents one observation.
- Response variable y , specified as an n -by-1 vector, where n is the number of observations. Each entry in y is the response for the corresponding row of X .

deepGLMfit

Description: deepGLMfit fits a deepGLM model using the design matrix X and response vector y, and returns an output structure **mdl** to make prediction on a test data. By default, if 'distribution' option is not specified, deepGLMfit will treat response variable y as normal distributed variable.

Syntax: mdl = deepGLMfit(X,y,
 dist = 'normal',
 initialize = 'adaptive',
 isotropic = False,
 seed = np.nan,
 nval = 0.2,
 verbose = 10,
 muTau = np.nan,
 lowerbound = True,
 windowSize = 100,
 network = [10,10],
 lrate = 0.01,
 S = 10,
 batchsize = 5000,
 epoch = 1000,
 tau = 10000,
 patience = 100,
 momentum = 0.6,
 Xval = [],
 yval = [],
 icept = True)

Input:

distr: Name of the distribution of the response, chosen from the following:

'normal'	Normal distribution (for continuous response)
'binomial'	Binomial distribution (for binary response)
'poisson'	Poisson distribution (for counting response)

Example	Distr = 'binomial'
Default	'normal'
Data type	single string

Network: Neuron Network structure for deepGLM. In the current version, deepGLM supports only 1 node for the output layer, users just need to provide a structure for hidden layers in an array where each element in the array is the number of nodes in the corresponding hidden layer.

Example	'Network', [10,10,10] → deepGLM uses 3 hidden layers with 10 nodes for each.
Default	[10,10]
Data type	array positive integer

lr: The fix learning rate that is used for training. If the learning rate is too small, training will take a long time, but if it is too high, the training is likely to get stuck at a suboptimal result.

Example	0.001
Default	0.01
Data type	single double

momentum: Momentum weight for stochastic gradient ascend. The momentum determines the contribution of the gradient step from the previous iteration to the current iteration of training. It must be a value between 0 and 1, where 0 will give no contribution from the previous step, and 1 will give a maximal contribution from the previous step. Must be between 0 and 1.

Example	0.9
Default	0.6
Data type	single double (0 to 1)

batchsize: The size of the mini-batch used for each training iteration. For deepGLM, batch size should be a large number (e.g. 5000, 10000) compared to batch size in deep learning literature (e.g. 128, 256). Must be a positive integer equal or smaller than number of observations of training data

Example	128
Default	5000
Data type	single positive integer

epoch: The maximum number of epochs that will be used for training. An epoch is defined as the number of iterations needed for optimization algorithm to scan entire training data. Must be a positive integer.

Example	1000
Default	100
Data type	single positive integer

patience: Number of consecutive times that the validation loss is allowed to be larger than or equal to the previously smallest loss before network training is stopped, used as an early stopping criterion. Must be a positive integer.

Example	20
Default	100
Data type	single positive integer

tau: Down-scaling factor that is applied to the learning rate every time a certain number of iterations has passed. Must be a positive integer

Example	100
Default	500
Data type	single positive integer

S: The number of samples needed for Monte Carlo approximation of gradient of lower bound. Must be an positive integer

Example	20
Default	10
Data type	single positive integer

windowSize: Size of moving average window that used to smooth the VB lowerbound. Must be an positive integer.

Example	50
Default	200
Data type	single positive integer

icept: Set true (default) to add a column of 1 to predictor observation X matrix (play the role as intercept). If the data have already included the first '1' column, set 'Intercept' to false.

Example	False
Default	True
Data type	single logical

verbose: Number of iterations that information on training progress will be printed to the command window each time. Set to 0 to disable this options.

Example	10 → Display training progress after each 5 iterations.
Default	0
Data type	single integer

isotropic: Set to true if you want to use Isotropic structure on Sigma (Variational Covariance matrix). By default, deepGLM uses Diagonal structure to factorize Sigma

Example	True
Default	False
Data type	single logical

seed: Seeds the random number generator using the nonnegative integer. Must be a nonnegative integer.

Example	500
Default	NaN
Data type	single nonnegative integer

Output: deepGLMfit returns a structure including all input setting by users and output training information with a combination of the following fields:

out.weights	Estimated Mean of weights of Deep Neuron Network from input layer to the last hidden layer. This will be used to quickly doing point estimation for new data.
out.beta	Estimated Mean of weights from the last hidden layer to output layer. This will be used to quickly doing point estimation for new data.
out.shrinkage	Matrix storing estimated values of group Lasso coefficients during training phase.
out.lb	A vector storing values of VB lowerbound calculated in each iteration. In the training algorithm, the convergence of lowerbound is used as early stopping rule.
out.lbBar	A vector storing values of VB lowerbound after smoothed by a moving average window.
out.vbMU	Estimated Mean of Variational Distribution. Used to do prediction interval estimation for new data.

out.vbSIGMA	Estimated Covariance matrix of Variational Distribution. Used to do prediction interval estimation for new data.
out.b	Estimated factor loading vector. Used to calculate the estimated Covariance matrix of Variational Distribution vbSIGMA
out.c	Estimated vector of idiosyncratic noise standard deviation. Used to calculate the estimated Covariance matrix of Variational Distribution vbSIGMA
out.sigma2Alpha	Estimated shape parameter of Invert Gamma Variation Approximation for variance of noise
out.sigma2Beta	Estimated scale parameter of Invert Gamma Variation Approximation for variance of noise
out.sigma2Mean	Estimated mean of covariance of noise in each iteration during training phase
out.nparams	Total number of parameters in deepGLM have to be trained.
out.indexTrack	A vector to keep track the indexes of all weights in Deep Neuron Network. Used to reconstruct Deep Neuron Network from vbMU when doing prediction on new data.
out.CPU	Total amount of training time (in second)

deepGLMpredict

Description: Predict responses for new data X_{test} using trained deepGLM structure mdl (output from deepGLMfit)

Syntax: deepGLMpredict(mdl,X, y=[], alpha=0, Nsample=1000, intercept=True)

Input:

y: Specify column of test responses. If this option is specified with true response column of new observations, deepGLMpredict will return prediction scores (PPS, MSE or Classification Rate) using true responses column vector y.

alpha: Return prediction interval estimation for observations in test data X_{test} . By default, this predictive interval capability is disable ('Interval' is 0). Must be an positive number.

Example 1 → Return one standard deviation predictive interval for new observations.

Default 0

Data type single | double

Nsample: Number of samples generated from posterior distribution of model parameters used to make prediction interval estimation for test data. Must be a positive integer

Example 200

Default 1000

Data type single | positive integer

Output: deepGLMpredict returns a structure of prediction results on test data with a combination of the following fields:

yhat	Vector of point estimation for Xtest. The length of yhat is equal to the number of observation of Xtest.
yNN	Vector of output values when pass Xtest through trained Deep Neuron Network. The length of yNN is the number of observation in Xtest. If responses are normal, yNN is equal to yhat.
yProb	Vector of estimated probabilities assigned for each observation in Xtest belong to class 1. This in only available for binary response data.
pps	Partial Predictive Score loss of trained deepGLM on Xtest. This is only available when ytest is specified.
mse	Mean Square Error of trained deepGLM on Xtest. This is only available for Normal or Poisson responses and ytest is specified.
accuracy	Classification rate of trained deepGLM classifier on Xtest. This is only available for binary responses and ytest is specified.
interval	Prediction interval estimation of trained deepGLM on Xtest. Interval is a two columns matrix where each row is an prediction interval of the corresponding observation of Xtest. This is only available when 'Interval' option is specified with a non-zero value.

deepGLMplot

Description: deepGLMplot plots data specified in data according the type specified by type.

Syntax: deepGLMplot(type, Pred,
 TextTitle=", labelX=", labelY=",
 linewidth=2, color='red', style='shade',
 npoint=50, order='ascend', y=[],legendText={}):

Input:

type	<ul style="list-style-type: none"> • 'Shrinkage' Plot stored values of group Lasso coefficient during training phase. If this type is specify, data is the output structure mdl from deepGLMfit or user can manually extract mdl.out.shrinkage field from mdl an use as input argument data. • 'Interval' Plot prediction interval estimation for test data. If this type is specified, data is the output structure Pred from deepGLMpredict. • 'ROC' Plot ROC curve for prediction from binary test data. If this type is specify, data is a matrix where the first column is a vector of target responses and the second column is the predicted vector yProb extract from output structure Pred of deepGLMpredict. If you want to plot different ROC curves, add more probability columns to data. - NOT IMPLEMENTED YET
npoint	This option only available when 'type' is Interval. 'Nsample' specifies

	number of test observations randomly selected from test data to plot prediction intervals.
TextTitle	Title of the figure
labelX	Label of X axis
labelY	Label of Y axis
linewidth	Line width of the plots
legendText	Legend creates a legend with descriptive labels for each plotted data series

Practical Recommendation

To make it easy to use deepGLM, we have already set the default settings for deepGLM that users can use in the first attempt of training a deepGLM model. However, these default setting can be inappropriate for some dataset that need a more sophisticated tuning of hyperparameters to effectively train a deepGLM model. Training deepGLM is hard because there are a lot of possible combinations of hyperparameters can be fed into deepGLMfit. There is always a trade-off between reducing training time and improving predictive performance of the trained model. The following recommendations give some heuristic guidelines when users want to improve the performance of deepGLM in each of the case:

1. Improving predictive performance of deepGLM model

a. Smoothing lowerbound

- DeepGLM uses the lowerbound of KL divergence between variational distribution and true posterior distribution of model parameters to make early stopping criterion and model selection. For this reason, we have to make sure that the model
- If the lowerbound does not have increasing trend or significantly fluctuate, then we can modify these hyperparameters to make a smooth lowerbound:
 - Increase batch size by modifying *batchSize* option of deepGLMfit. Increasing batch size helps reducing the variance of Monte Carlo estimation of loglikelihood contribution of lowerbound.
 - Increase moving average window size by modifying *windowSize* option of deepGLMfit. Increasing window size helps smoothing fluctuated lowerbound.
 - Reducing learning rate by modifying *lr* option. This helps lowerbound not ‘overshoot’ when it moves closely to saturated region.
 - Increase *S* option of deepGLMfit (e.g. $S=20$). This helps to reduce the variance of MC estimation of log-likelihood contribution of lowerbound

b. Try various neuron network structures

- Modify option *network* of deepGLMfit
- An inappropriate structure may cause high bias on test data.
- Start with a simple network (e.g. $\text{network} = [5,5]$) then increasing the network complexity until the predictive performance gets worse on test data.
- If the network is too complicated, it will take longer time for training. Additionally, increasing network complexity (adding more layers) does not ensure an improvement in predictive performance of deepGLM

c. Decreasing learning rate

- Decrease *lr* option of deepGLMfit
- A too small learning rate will slow down training process.
- If the total number of iterations needed for training is smaller than *tau*, then reducing *tau* could help to decrease learning rate.

d. Increase number of training iterations

- Increase *patience* option of deepGLMfit. This is to relax the early stopping criteria.
- Increase *epoch* option of deepGLMfit. This is to relax the early stopping criteria.

- e. Others
 - Try different valid values of *momentum*, *batchSize* options may help to improve prediction accuracy on test data.
 - Increase *S* option of deepGLMfit (e.g. $S = 20$). A small value of *S* may cause high variation in the Monte Carlo estimation of gradient of lower bound.
2. Reduce training time
- a. Increase learning rate

Increase *lr* option of deepGLMfit (e.g. *lr* = 0.1)

If learning is too high, training algorithm may be failed to converge
 - b. Decrease number of training iterations

Decrease *patience* option of deepGLMfit (e.g. *patience* = 10).

Increase *epoch* option of deepGLMfit.
 - c. Simplify Neuron Network structure
 - Set simpler structure for *network* option of deepGLMfit.
 - A too simple Neuron Network structure does not imply that deepGLM will have worse performance. The complexity of Neuron structure depends on the complexity of training data.
 - d. Use Isotropic structure for very complex neuron network structure
 - Enable *isotropic* option of deepGLMfit to true (e.g. *isotropic* = True)
 - e. Others

Disable *monitor* option of deepGLMfit to *False* (e.g. *monitor* = False)

Decrease *S* option of deepGLMfit (e.g. $S = 5$). A small number of *S* may cause high variation in the Monte Carlo estimation of gradient of lower bound.